

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: import os, types
import pandas as pd
from botocore.client import Config
import ibm_boto3

def __iter__(self): return 0

#@hidden_cell
# The following code accesses a file in your IBM Cloud Object Storage. It includes your credentials.
# You might want to remove those credentials before you share the notebook.
cos_client = ibm_boto3.client(service_name='s3',
                              ibm_api_key_id='Uq1L_fr0mGem56BJ158Ro0C5ks-jybAjXNehBzPseAZ1',
                              ibm_auth_endpoint='https://iam.cloud.ibm.com/oidc/token',
                              config=Config(signature_version='oauth'),
                              endpoint_url='https://s3.private.us.cloud-object-storage.appdomain.cloud')

bucket = 'carresalevalue-donotdelete-pr-2lv9juqbt5eqf'
object_key = 'autos.csv'

body = cos_client.get_object(Bucket=bucket, Key=object_key)['Body']
# add missing __iter__ method, so pandas accepts body as file-like object
if not hasattr(body, "__iter__"): body.__iter__ = types.MethodType(__iter__, body)

df = pd.read_csv(body)
df.head()
```

```
Out[2]:
```

	dateCrawled	name	seller	offerType	price	abtest	vehicleType	yearOfRegistration	gearbox	powerPS
0	2016-03-24 11:52:17	Golf_3_1.6	privat	Angebot	480	test	NaN	1993	manuell	0
1	2016-03-24 10:58:45	A5_Sportback_2.7_Tdi	privat	Angebot	18300	test	coupe	2011	manuell	190
2	2016-03-14 12:52:21	Jeep_Grand_Cherokee_Overland	privat	Angebot	9800	test	suv	2004	automatik	163
3	2016-03-17 16:54:04	GOLF_4_1.4_3TURER	privat	Angebot	1500	test	kleinwagen	2001	manuell	75
4	2016-03-31 17:25:20	Skoda_Fabia_1.4_TDI_PD_Classic	privat	Angebot	3600	test	kleinwagen	2008	manuell	69

```
In [3]: # df = pd.read_csv('autos.csv')
```

```
In [4]: df.drop(columns= ['dateCrawled', 'name', 'seller', 'offerType', 'model', 'dateCreated',
                          'nrOfPictures', 'postalCode', 'lastSeen'], inplace = True)
```

```
In [5]: df
```

```
Out[5]:
```

	price	abtest	vehicleType	yearOfRegistration	gearbox	powerPS	kilometer	monthOfRegistration	fuelType	brand
0	480	test	NaN	1993	manuell	0	150000	0	benzin	volkswagen
1	18300	test	coupe	2011	manuell	190	125000	5	diesel	audi
2	9800	test	suv	2004	automatik	163	125000	8	diesel	jeep
3	1500	test	kleinwagen	2001	manuell	75	150000	6	benzin	volkswagen
4	3600	test	kleinwagen	2008	manuell	69	90000	7	diesel	skoda
...	...	...	...	...	...	...	...	...	...	...
371523	2200	test	NaN	2005	NaN	0	20000	1	NaN	sonstige_autos
371524	1199	test	cabrio	2000	automatik	101	125000	3	benzin	smart
371525	9200	test	bus	1996	manuell	102	150000	3	diesel	volkswagen
371526	3400	test	kombi	2002	manuell	100	150000	6	diesel	volkswagen
371527	28990	control	limousine	2013	manuell	320	50000	8	benzin	bmw

371528 rows x 11 columns

```
In [6]: # Removing Missing Values
df['vehicleType'].fillna(df['vehicleType'].mode()[0], inplace = True)
df['gearbox'].fillna(df['gearbox'].mode()[0], inplace = True)
df['fuelType'].fillna(df['fuelType'].mode()[0], inplace = True)
df['notRepairedDamage'].fillna(df['notRepairedDamage'].mode()[0], inplace = True)
```

```
In [7]: df.isna().sum()
```

```
Out[7]: price      0
abtest      0
vehicleType  0
yearOfRegistration  0
gearbox      0
powerPS      0
kilometer    0
monthOfRegistration  0
fuelType     0
brand        0
notRepairedDamage  0
dtype: int64
```

```
In [8]: # Checking for Duplicates
df.duplicated().sum()
```

```
Out[8]: 24211
```

```
In [9]: df = df.drop_duplicates()
```

```
In [10]: df.duplicated().sum()

Out[10]: 0

In [11]: df
```

	price	abtest	vehicleType	yearOfRegistration	gearbox	powerPS	kilometer	monthOfRegistration	fuelType	brand
0	480	test	limousine	1993	manuell	0	150000	0	benzin	volkswagen
1	18300	test	coupe	2011	manuell	190	125000	5	diesel	audi
2	9800	test	suv	2004	automatik	163	125000	8	diesel	jeep
3	1500	test	kleinwagen	2001	manuell	75	150000	6	benzin	volkswagen
4	3600	test	kleinwagen	2008	manuell	69	90000	7	diesel	skoda
...	...	...	...	...	...	...	...	...	...	...
371523	2200	test	limousine	2005	manuell	0	20000	1	benzin	sonstige_autos
371524	1199	test	cabrio	2000	automatik	101	125000	3	benzin	smart
371525	9200	test	bus	1996	manuell	102	150000	3	diesel	volkswagen
371526	3400	test	kombi	2002	manuell	100	150000	6	diesel	volkswagen
371527	28990	control	limousine	2013	manuell	320	50000	8	benzin	bmw

347317 rows x 11 columns

```
In [12]: # Dropping outliers in price
a = df[df['price'] > 50000].index
df.drop(a, inplace = True)

In [13]: # Dropping outliers in yearOfRegistration
a = df[df['yearOfRegistration'] > 2019].index
df.drop(a, inplace = True)
a = df[df['yearOfRegistration'] < 1985].index
df.drop(a, inplace = True)

In [14]: # Dropping outliers in powerPS
a = df[df['powerPS'] > 300].index
df.drop(a, inplace = True)
a = df[df['powerPS'] < 40].index
df.drop(a, inplace = True)

In [15]: # Label Encoding
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
df['abtest'] = le.fit_transform(df['abtest'])
df['vehicleType'] = le.fit_transform(df['vehicleType'])
df['gearbox'] = le.fit_transform(df['gearbox'])
df['fuelType'] = le.fit_transform(df['fuelType'])
df['brand'] = le.fit_transform(df['brand'])
df['notRepairedDamage'] = le.fit_transform(df['notRepairedDamage'])

In [16]: df
```

	price	abtest	vehicleType	yearOfRegistration	gearbox	powerPS	kilometer	monthOfRegistration	fuelType	brand	notRepair
1	18300	1	3	2011	1	190	125000	5	3	1	
2	9800	1	7	2004	0	163	125000	8	3	14	
3	1500	1	4	2001	1	75	150000	6	1	38	
4	3600	1	4	2008	1	69	90000	7	3	31	
5	650	1	6	1995	1	102	150000	10	1	2	
...	...	...	...	...	...	...	...	...	...	...	...
371519	5250	0	6	2016	0	150	150000	12	1	0	
371520	3200	0	6	2004	1	225	150000	5	1	30	
371524	1199	1	2	2000	0	101	125000	3	1	32	
371525	9200	1	1	1996	1	102	150000	3	3	38	
371526	3400	1	5	2002	1	100	150000	6	3	38	

298133 rows x 11 columns

```
In [17]: # Splitting x and y variables
x = df.drop(columns = 'price')
y = df['price']

In [18]: # Splitting into test and train
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=0)
```

## Building Models

```
In [18]: # Linear Regression

In [19]: from sklearn.linear_model import LinearRegression
lr = LinearRegression()
lr.fit(x_train, y_train)

Out[19]: LinearRegression()

In [20]: # Lasso Regression

In [21]: from sklearn.linear_model import Lasso
lasso = Lasso(alpha=0.01, normalize=True)
lasso.fit(x_train, y_train)
```

/opt/conda/envs/Python-3.9/lib/python3.9/site-packages/sklearn/linear\_model/\_base.py:141: FutureWarning: 'normalize' was deprecated in version 1.0 and will be removed in 1.2. If you wish to scale the data, use Pipeline with a StandardScaler in a preprocessing stage. To reproduce the previous behavior:

```
Out[21]: Lasso(alpha=0.01, normalize=True)
```

```
In [22]: # Ridge Regression
```

```
In [23]: from sklearn.linear_model import Ridge
ridge = Ridge(alpha=0.01, normalize=True)
ridge.fit(x_train, y_train)
```

/opt/conda/envs/Python-3.9/lib/python3.9/site-packages/sklearn/linear\_model/\_base.py:141: FutureWarning: 'normalize' was deprecated in version 1.0 and will be removed in 1.2.  
If you wish to scale the data, use Pipeline with a StandardScaler in a preprocessing stage. To reproduce the previous behavior:

```
from sklearn.pipeline import make_pipeline
```

```
model = make_pipeline(StandardScaler(with_mean=False), Ridge())
```

If you wish to pass a sample\_weight parameter, you need to pass it as a fit parameter to each step of the pipeline as follows:

```
kwargs = {s[0] + '._sample_weight': sample_weight for s in model.steps}
model.fit(X, y, **kwargs)
```

```
Set parameter alpha to: original_alpha * n_samples.
warnings.warn(
```

```
Out[23]: Ridge(alpha=0.01, normalize=True)
```

```
In [24]: # Decision Tree
```

```
In [25]: from sklearn.tree import DecisionTreeRegressor
DT = DecisionTreeRegressor()
DT.fit(x_train, y_train)
```

```
Out[25]: DecisionTreeRegressor()
```

```
In [26]: # KNN
```

```
In [27]: from sklearn.neighbors import KNeighborsRegressor
knn = KNeighborsRegressor()
knn.fit(x_train, y_train)
```

```
Out[27]: KNeighborsRegressor()
```

```
In [28]: # Random Forest
```

```
In [19]: from sklearn.ensemble import RandomForestRegressor
RF = RandomForestRegressor()
RF.fit(x_train, y_train)
```

```
Out[19]: RandomForestRegressor()
```

```
In [30]: # Linear Regression
acc_lr = lr.score(x_test, y_test)*100
```

```
In [33]: from sklearn.metrics import mean_squared_error
np.sqrt(mean_squared_error(y_test, lr.predict(x_test)))
```

```
Out[33]: 3956.5815636035836
```

```
In [34]: # Lasso Regression
acc_lasso = lasso.score(x_test, y_test)*100
```

```
In [35]: np.sqrt(mean_squared_error(y_test, lasso.predict(x_test)))
```

```
Out[35]: 3956.664250218968
```

```
In [36]: # Ridge Regression
acc_ridge = ridge.score(x_test, y_test)*100
```

```
In [37]: np.sqrt(mean_squared_error(y_test, ridge.predict(x_test)))
```

```
Out[37]: 3956.8355512022918
```

```
In [38]: # K Nearest Neighbour
acc_knn = knn.score(x_test, y_test)*100
```

```
In [39]: np.sqrt(mean_squared_error(y_test, knn.predict(x_test)))
```

```
Out[39]: 2675.1305595164295
```

```
In [40]: # Decision Tree
acc_dt = DT.score(x_test, y_test)*100
```

```
In [41]: np.sqrt(mean_squared_error(y_test, DT.predict(x_test)))
```

```
Out[41]: 2936.6216848726476
```

```
In [42]: # Random Forest
acc_rf = RF.score(x_test, y_test)*100
```

```
In [43]: np.sqrt(mean_squared_error(y_test, RF.predict(x_test)))
```

```
Out[43]: 2247.585702127866
```

```
In [44]: models = pd.DataFrame(
    {'Model': ['Linear Regression', 'Lasso Regression', 'Ridge Regression', 'KNN',
               'Decision Tree', 'Random Forest'],
     'Score': [acc_lr, acc_lasso, acc_ridge, acc_knn, acc_dt, acc_rf]}
```

ite-packages (from packaging->ibm-watson-machine-learning) (3.0.4)  
Note: you may need to restart the kernel to use updated packages.

```
In [25]: from ibm_watson_machine_learning import APIClient
import json
wml_credentials = {'apikey': 'gltklKa55mB-mBztAhCCv8gMemWgixeSSsD4qW0Smqr9',
                  'url': 'https://us-south.ml.cloud.ibm.com'}
wml_client = APIClient(wml_credentials)

In [20]: # def guid_from_space_name(client, space_name):
#         space = client.spaces.get_details()
#         return(next(iter for item in space['resources'] if item['entity']['name'] == space_name)['metao

In [21]: # space_uid = guid_from_space_name(client, 'model1')
# print('space_uid = ' + space_uid)

In [26]: wml_client.spaces.list()

Note: 'limit' is not provided. Only first 50 records will be displayed if the number of records excee
d 50
-----
ID                NAME      CREATED
cb3b5444-ea5c-4352-914d-cf7569dd8c9f  model1    2022-11-12T06:55:56.978Z
-----

In [27]: space_id = 'cb3b5444-ea5c-4352-914d-cf7569dd8c9f'

In [29]: wml_client.set.default_space(space_id)

Out[29]: 'SUCCESS'

In [42]: client.software_specifications.list()

-----
NAME                ASSET_ID      TYPE
default_py3.6       0062b8c9-8b7d-44a0-a9b9-46c416adcbd9  base
kernel-spark3.2-scala2.12  020d69ce-7ac1-5e68-ac1a-31189867356a  base
pytorch-onnx_1.3-py3.7-edt  069ea134-3346-5748-b513-49120e15d288  base
scikit-learn_0.20-py3.6    09c5a1d0-9c1e-4473-a344-eb7b665ff687  base
spark-mllib_3.0-scala_2.12  09f4cff0-90a7-5899-b9ed-1ef348aebdee  base
pytorch-onnx_rt22.1-py3.9   0b848dd4-e681-5599-be41-b5f6fcccc6471  base
ai-function_0.1-py3.6       0cdb0f1e-5376-4f4d-92dd-da3b69aa9bda  base
shiny-r3.6                 0e6e79df-875e-4fd4-8ae9-62dcc2148306  base
tensorflow_2.4-py3.7-horovod  1092590a-307d-563d-9b62-4eb7d64b3f22  base
pytorch_1.1-py3.6           10ac12d6-6b30-4ccd-8392-3e922c096a92  base
tensorflow_1.15-py3.6-ddl    111e41b3-de2d-5422-a4d6-bf776828c4b7  base
autoai-kb_rt22.2-py3.10     125b6d9a-5b1f-5e8d-972a-b251688ccf40  base
runtime-22.1-py3.9          12b83a17-24d8-5082-900f-0ab31fbfd3cb  base
scikit-learn_0.22-py3.6     154010fa-5b3b-4ac1-82af-4d5ee5abbc85  base
default_r3.6               1b70aec3-ab34-4b87-8aa0-a4a3c8296a36  base
pytorch-onnx_1.3-py3.6       1bc6029a-cc97-56da-b8e0-39c3880dbbe7  base
kernel-spark3.3-r3.6         1c9e5454-f216-59dd-a20e-474a5cdf5988  base
pytorch-onnx_rt22.1-py3.9-edt  1d362186-7ad5-5b59-8b6c-9d0880bde37f  base
tensorflow_2.1-py3.6         1eb25b84-d6ed-5dde-b6a5-3fbdff1665666  base
spark-mllib_3.2             20047f72-0a98-58c7-9ff5-a77b012eb8f5  base
tensorflow_2.4-py3.8-horovod  217c16f6-178f-56bf-824a-b19f20564c49  base
runtime-22.1-py3.9-cuda     26215f05-08c3-5a41-a1b0-da66306ce658  base
do_py3.8                   295addb5-9ef9-547e-9bf4-92ae3563e720  base
autoai-ts_3.8-py3.8         2aa0c932-798f-5ae9-abd6-15e0c2402fb5  base
tensorflow_1.15-py3.6       2b73a275-7cbf-420b-a912-eae7f436e0bc  base
kernel-spark3.3-py3.9       2b7961e2-e3b1-5a8c-a491-482c8368839a  base
pytorch_1.2-py3.6           2c8ef57d-2687-4b7d-acce-01f94976dac1  base
spark-mllib_2.3             2e51f700-bca0-4b0d-88dc-5c6791338875  base
pytorch-onnx_1.1-py3.6-edt   32983cea-3f32-4400-8965-dde874a8d67e  base
spark-mllib_3.0-py37        36507ebe-8770-55ba-ab2a-eafe787600e9  base
spark-mllib_2.4             390d21f8-e58b-4fac-9c55-d7ceda621326  base
autoai-ts_rt22.2-py3.10     396b2e83-0953-5b86-9a55-7ce1628a406f  base
xgboost_0.82-py3.6          39e31acd-5f30-41dc-ae44-60233c80306e  base
pytorch-onnx_1.2-py3.6-edt   40589d0e-7019-4e28-8daa-fb03b6f4fe12  base
pytorch-onnx_rt22.2-py3.10   40e73f55-783a-5535-b3fa-0c8b94291431  base
default_r36py38             41c247d3-45f8-5a71-b065-8580229facf0  base
autoai-ts_rt22.1-py3.9       4269d26e-07ba-5d40-8f66-2d495b0c71f7  base
autoai-obm_3.0              42b92e18-d9ab-567f-988a-4240ba1ed5f7  base
pmml-3.0_4.3                493bcb95-16f1-5bc5-bee8-81baf80e9c7  base
spark-mllib_2.4-r_3.6       49403dff-92e9-4c87-a3d7-a42d0021c095  base
xgboost_0.90-py3.6          4ff8d6c2-1343-4c18-85e1-689c965304d3  base
pytorch-onnx_1.1-py3.6       50f95b2a-bc16-43bb-bc94-b0bed208c60b  base
autoai-ts_3.9-py3.8         52c57136-80fa-572e-8728-a5e7cbb42cde  base
spark-mllib_2.4-scala_2.11   55a70f99-7320-4be5-9fb9-9edb5a443af5  base
spark-mllib_3.0              5c1b0ca2-4977-5c2e-9439-ffd44ea8ffe9  base
autoai-obm_2.0              5c2e37fa-80b8-5e77-840f-d912469614ee  base
spss-modeler_18.1           5c3cad7e-507f-4b2a-a9a3-ab53a21dee8b  base
cuda-py3.8                  5d3232bf-c86b-5df4-a2cd-7bb870a1cd4e  base
autoai-kb_3.1-py3.7         632d4b22-10aa-5180-88f0-f52dfb6444d7  base
pytorch-onnx_1.7-py3.8      634d3cdc-b562-5bf9-a2d4-ea90a478456b  base
-----
Note: Only first 50 records were displayed. To display more use 'limit' parameter.

In [56]: import sklearn
sklearn.__version__

Out[56]: '1.0.2'

In [30]: model_name = 'rf'
deployment_name = 'rf'
demo_model = RF

In [31]: software_spec_uid = wml_client.software_specifications.get_id_by_name('runtime-22.1-py3.9')
software_spec_uid

Out[31]: '12b83a17-24d8-5082-900f-0ab31fbfd3cb'

In [32]: model_props = {
wml_client.repository.ModelMetaNames.NAME: model_name,
wml_client.repository.ModelMetaNames.TYPE: 'scikit-learn_1.0',
wml_client.repository.ModelMetaNames.SOFTWARE_SPEC_UID: software_spec_uid
}

In [33]: model_details = wml_client.repository.store_model(
model = demo_model,
meta_props=model_props,
training_data = x_train,
training_target = y_train)
```



autoai-obm_2.0	5c2e37fa-80b8-5e77-840f-d912469614ee	base
spss-modeler_18.1	5c3cad7e-507f-4b2a-a9a3-ab53a21dee8b	base
cuda-py3.8	5d3232bf-c86b-5df4-a2cd-7bb870a1cd4e	base
autoai-kb_3.1-py3.7	632d4b22-10aa-5180-88f0-f52dfb6444d7	base
pytorch-onnx_1.7-py3.8	634d3cdc-b562-5bf9-a2d4-ea90a478456b	base

Note: Only first 50 records were displayed. To display more use 'limit' parameter.

```
In [56]: import sklearn
sklearn.__version__
```

```
Out[56]: '1.0.2'
```

```
In [30]: model_name = 'rf'
deployment_name = 'rf'
demo_model = RF
```

```
In [31]: software_spec_uid = wml_client.software_specifications.get_id_by_name('runtime-22.1-py3.9')
software_spec_uid
```

```
Out[31]: '12b83a17-24d8-5082-900f-0ab31fbfd3cb'
```

```
In [32]: model_props = {
wml_client.repository.ModelMetaNames.NAME: model_name,
wml_client.repository.ModelMetaNames.TYPE: 'scikit-learn_1.0',
wml_client.repository.ModelMetaNames.SOFTWARE_SPEC_UID: software_spec_uid
}
```

```
In [33]: model_details = wml_client.repository.store_model(
model = demo_model,
meta_props=model_props,
training_data = x_train,
training_target = y_train)
```

```
In [34]: model_details
```

```
Out[34]: {'entity': {'hybrid_pipeline_software_specs': [],
'label_column': 'price',
'schemas': {'input': [{'fields': [{'name': 'abtest', 'type': 'int64'},
{'name': 'vehicleType', 'type': 'int64'},
{'name': 'yearOfRegistration', 'type': 'int64'},
{'name': 'gearbox', 'type': 'int64'},
{'name': 'powerPS', 'type': 'int64'},
{'name': 'kilometer', 'type': 'int64'},
{'name': 'monthOfRegistration', 'type': 'int64'},
{'name': 'fuelType', 'type': 'int64'},
{'name': 'brand', 'type': 'int64'},
{'name': 'notRepairedDamage', 'type': 'int64'}]},
'id': '1',
'type': 'struct'}]},
'output': [],
'software_spec': {'id': '12b83a17-24d8-5082-900f-0ab31fbfd3cb',
'name': 'runtime-22.1-py3.9',
'type': 'scikit-learn_1.0'},
'metadata': {'created_at': '2022-11-12T08:02:16.890Z',
'id': 'dfaf6ba0-6752-4846-abb2-91f5c4ff27e0',
'modified_at': '2022-11-12T08:08:20.432Z',
'name': 'rf',
'owner': 'IBMId-6660020F0N',
'resource_key': '5d76cf1d-a018-474e-824d-11a308540d6c',
'space_id': 'cb3b5444-ea5c-4352-914d-cf7569dd8c9f',
'system': {'warnings': []}}}
```

```
In [35]: model_id = wml_client.repository.get_model_id(model_details)
```

```
In [36]: model_id
```

```
Out[36]: 'dfaf6ba0-6752-4846-abb2-91f5c4ff27e0'
```

```
In [37]: deployment_props = {
wml_client.deployments.ConfigurationMetaNames.NAME: deployment_name,
wml_client.deployments.ConfigurationMetaNames.ONLINE: {}
}
```

```
In [38]: deployment = wml_client.deployments.create(
artifact_uid = model_id,
meta_props = deployment_props)
```

#####

Synchronous deployment creation for uid: 'dfaf6ba0-6752-4846-abb2-91f5c4ff27e0' started

#####

initializing

Note: online\_url is deprecated and will be removed in a future release. Use serving\_urls instead.

....

ready

-----  
Successfully finished deployment creation, deployment\_uid='7a83cc3a-f415-4c17-9e47-da929f749bc5'

```
In [ ]:
```

[Give feedback](#)